

Namespace

- Organise resources in namespaces
- Virtual cluster inside cluster

Kubernetes provides 4 default namespaces

*Kubernetes-dashboard → only with minikube
Kube-system → do not create or modify in kube-system

- system processes are present
- Master & kubectl processes are present

Kube-public → publicly accessible data
configMap data, which contains cluster information

Kube-node-lease → heartbeats of nodes
each node has associated lease object in namespace
determines the availability of node

default → resources you create are located here

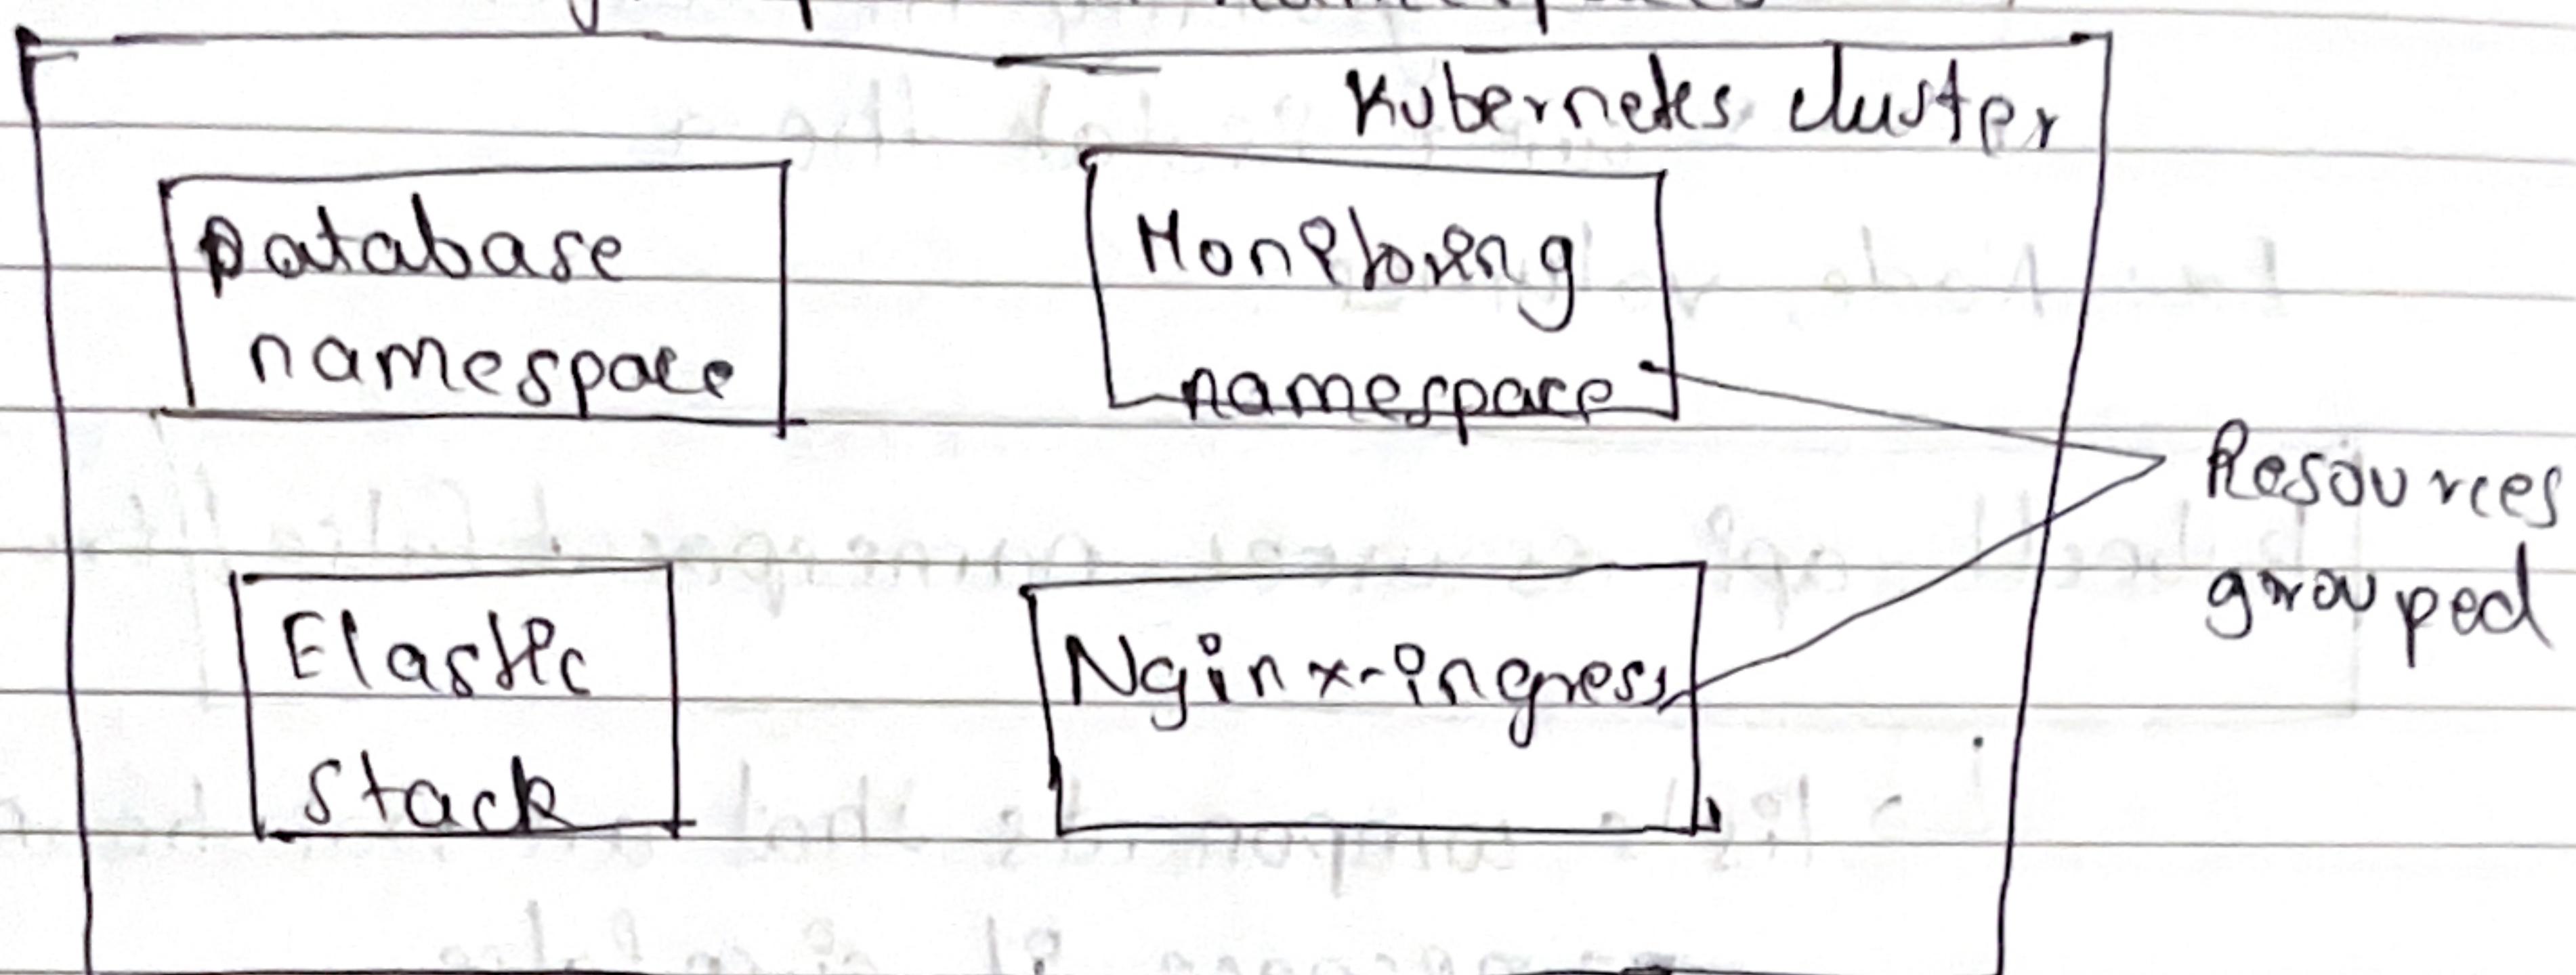
Note:

Create namespace with configuration file

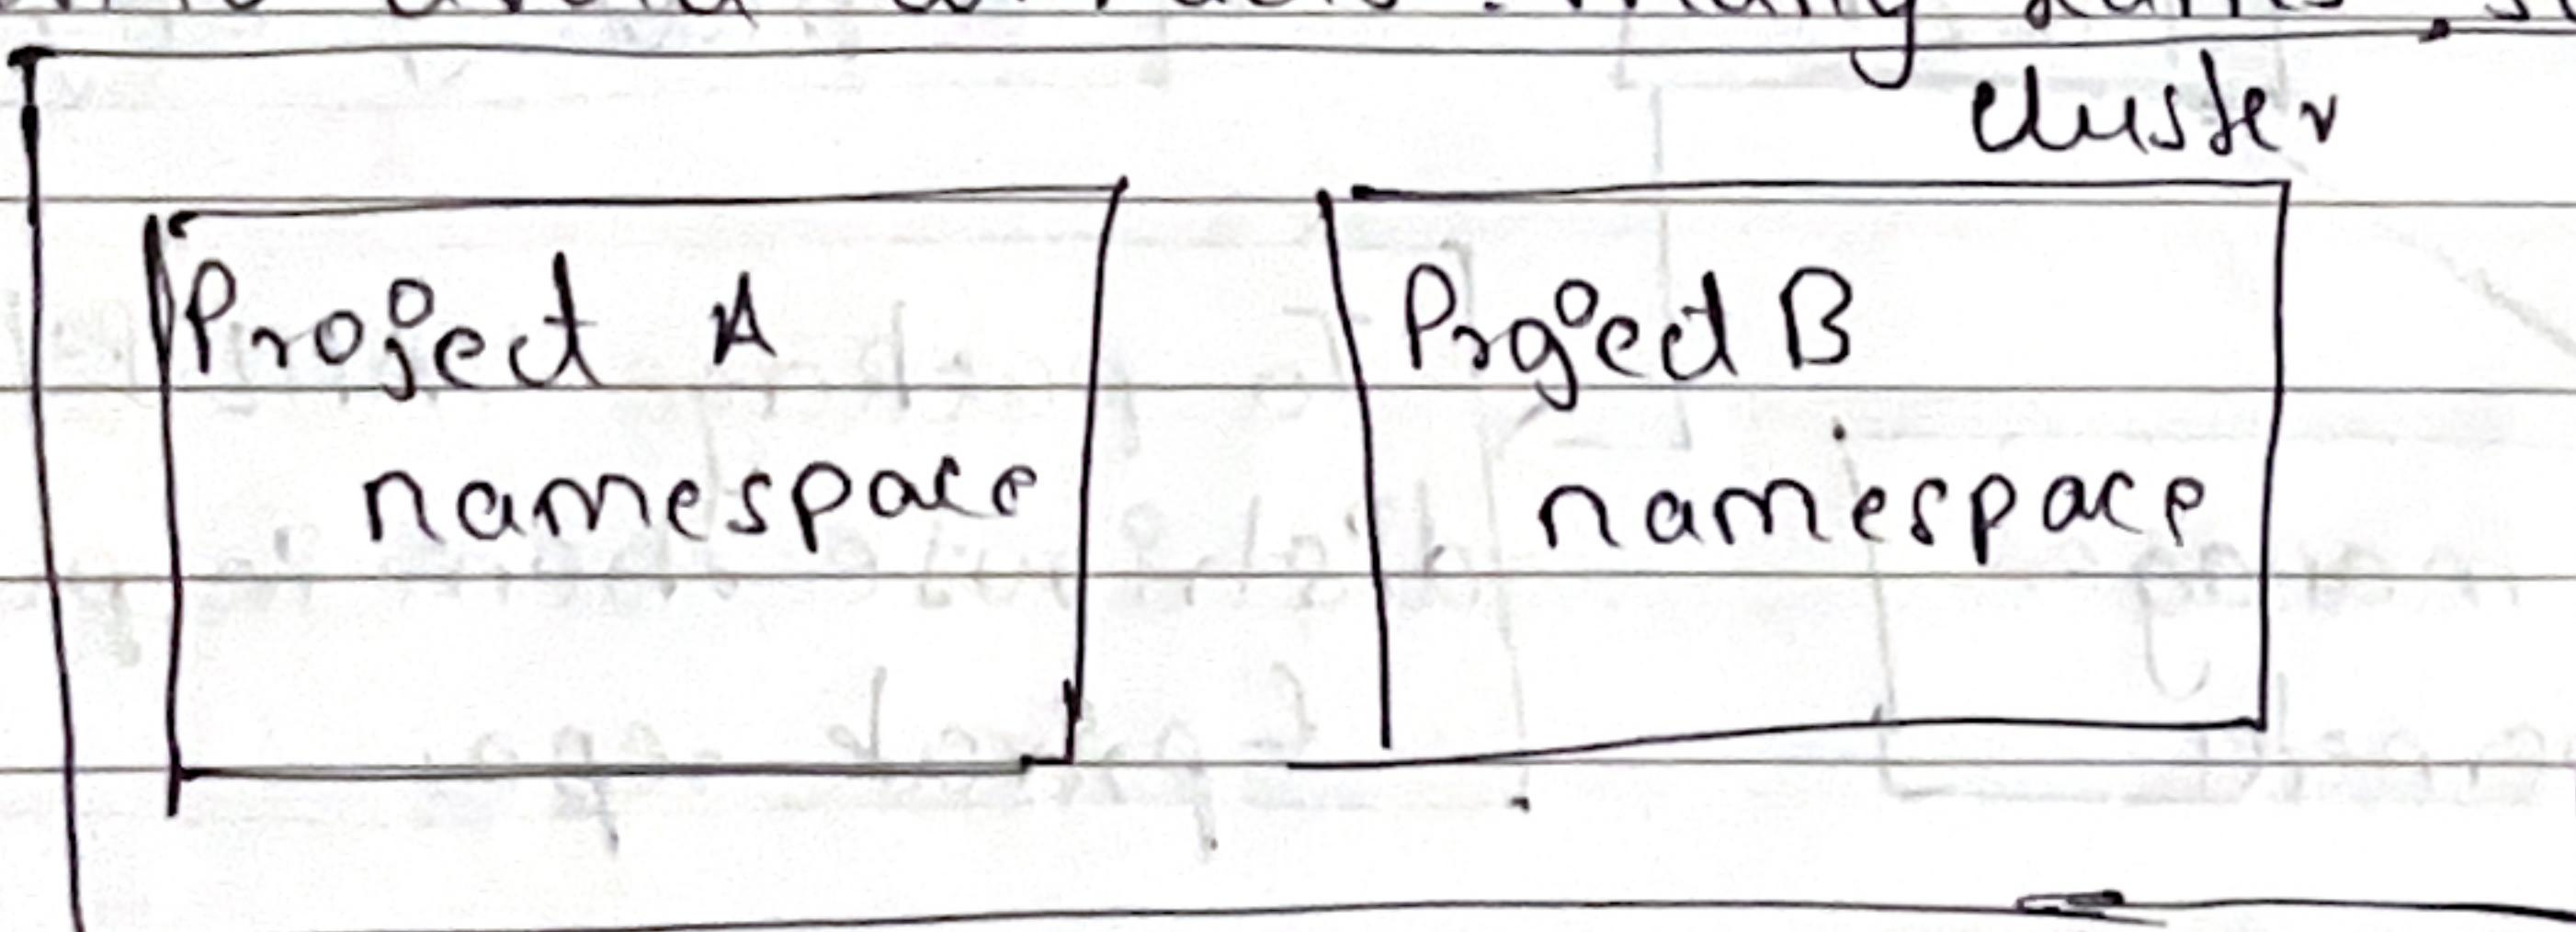
Officially: should not use for smaller projects

Why need for namespace?

- i) Resources grouped in namespaces



- ii) ~~Goal~~ To avoid conflicts: many teams, same application



- iii) Resource sharing: staging & development

- iv) Blue Green development like different versions of application

- v) Access & resource limits on namespace between teams

Note:

- vi) You can't access most resources from another namespace like you ~~should~~ can't use or access configMap or secret of other namespace.

Note:

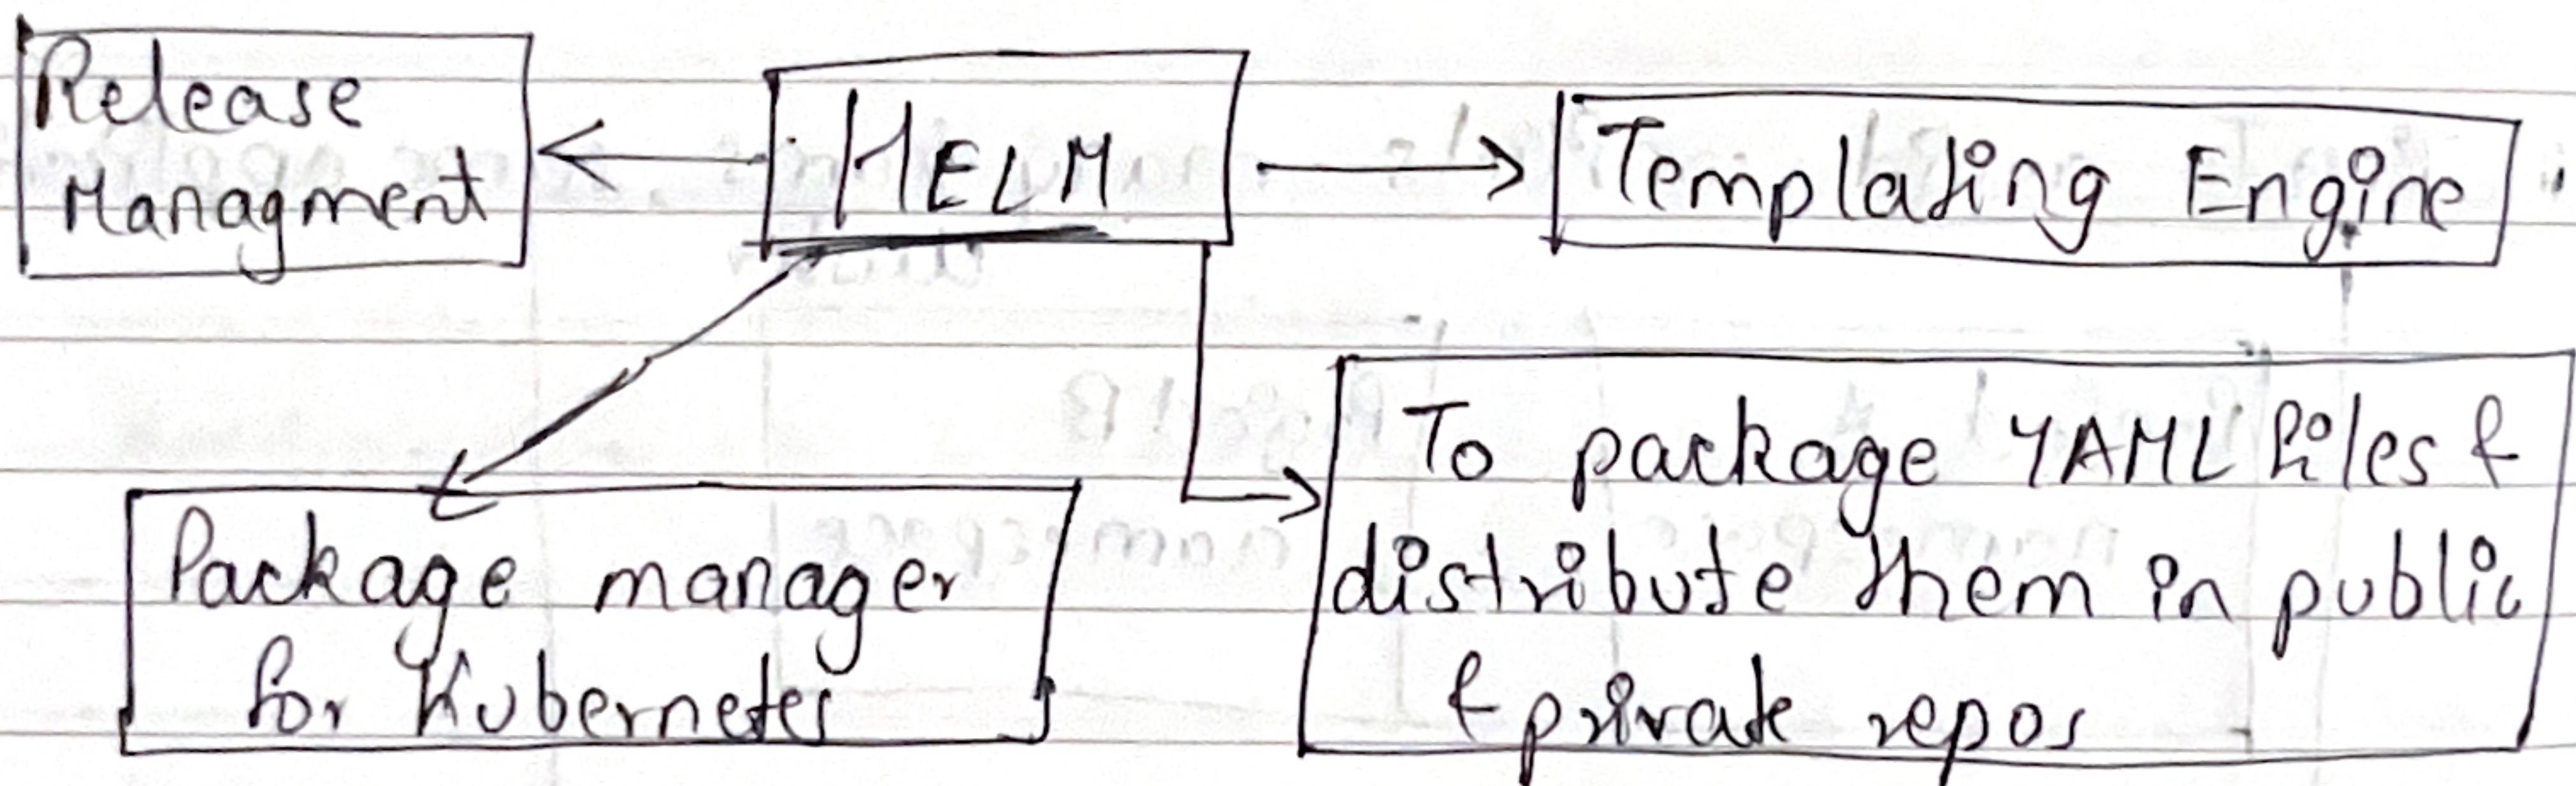
Some components cannot be created within namespace because

- → live globally on cluster
- → can't isolate them

Eg: Node, volumes

`kubectl api-resources --namespaced=false / true`

↳ lists components that are not bound to namespace if given false



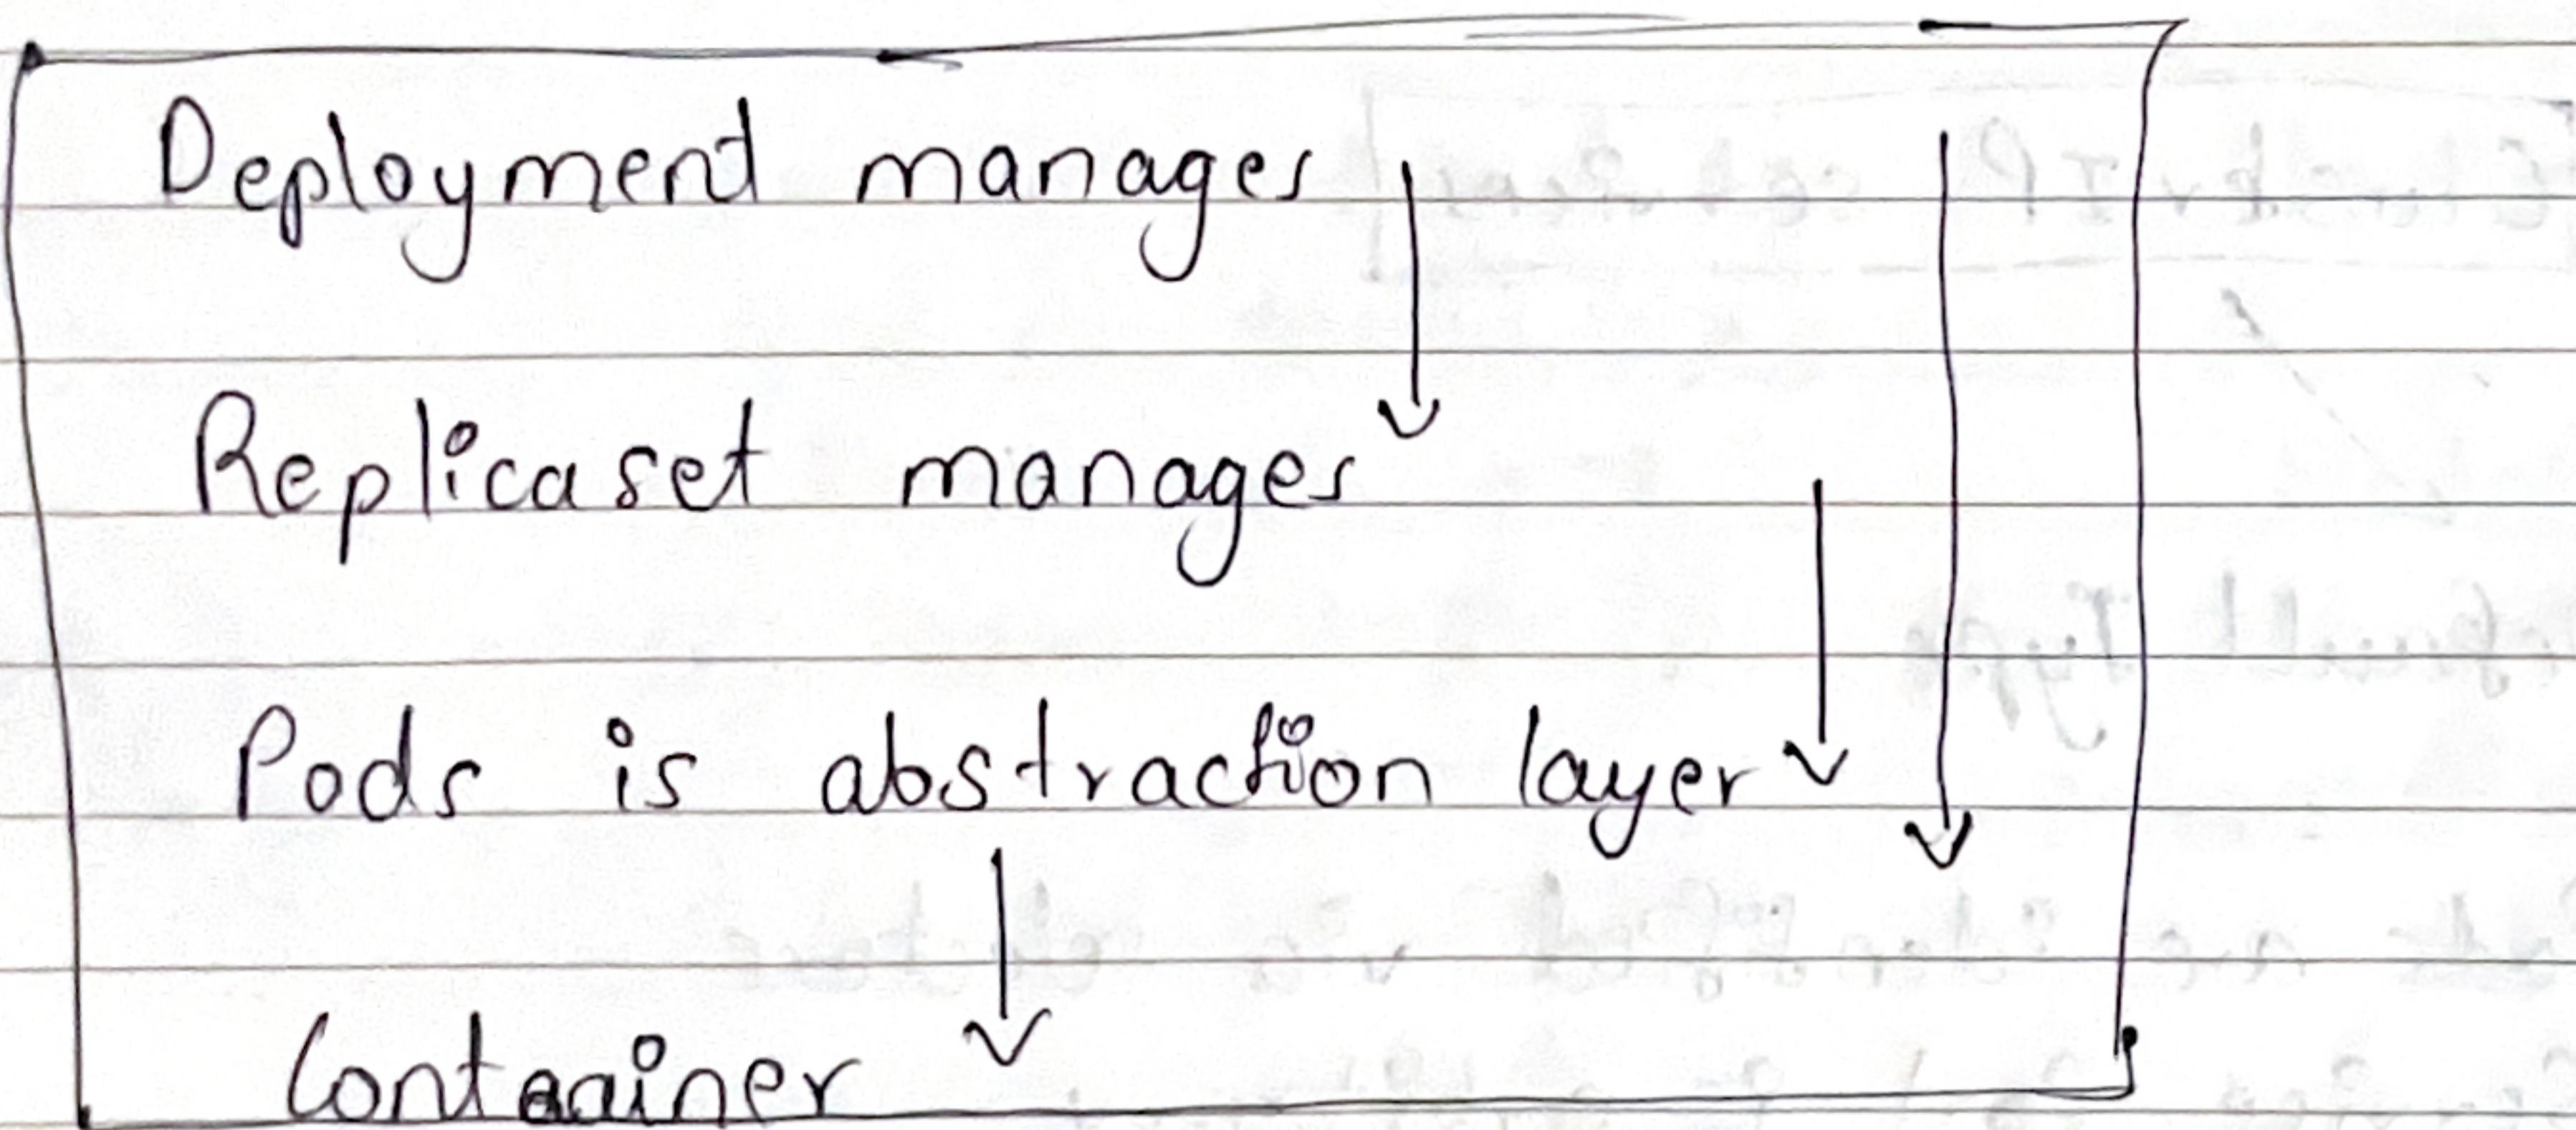
Helm charts

- Bundle of yaml files
- create your own Helm charts with helm
- Push them to helm repo
- download & use existing ones

YAML file → strict indentation

Each configuration file has 3 parts:

- i) Meta data
- Specification
- Status [automatically generated & added by Kubernetes]



Labels & selectors

↳ for connection establishment

Array / lists

Fruits:

- Orange
- Apple

Vegetables:

- Carrot
- Tomato

Dictionary / Map

Banana:

Calories: 105
Fat: 0.4 g

Grapes:

Calories: 62
Fat: 0.2 g

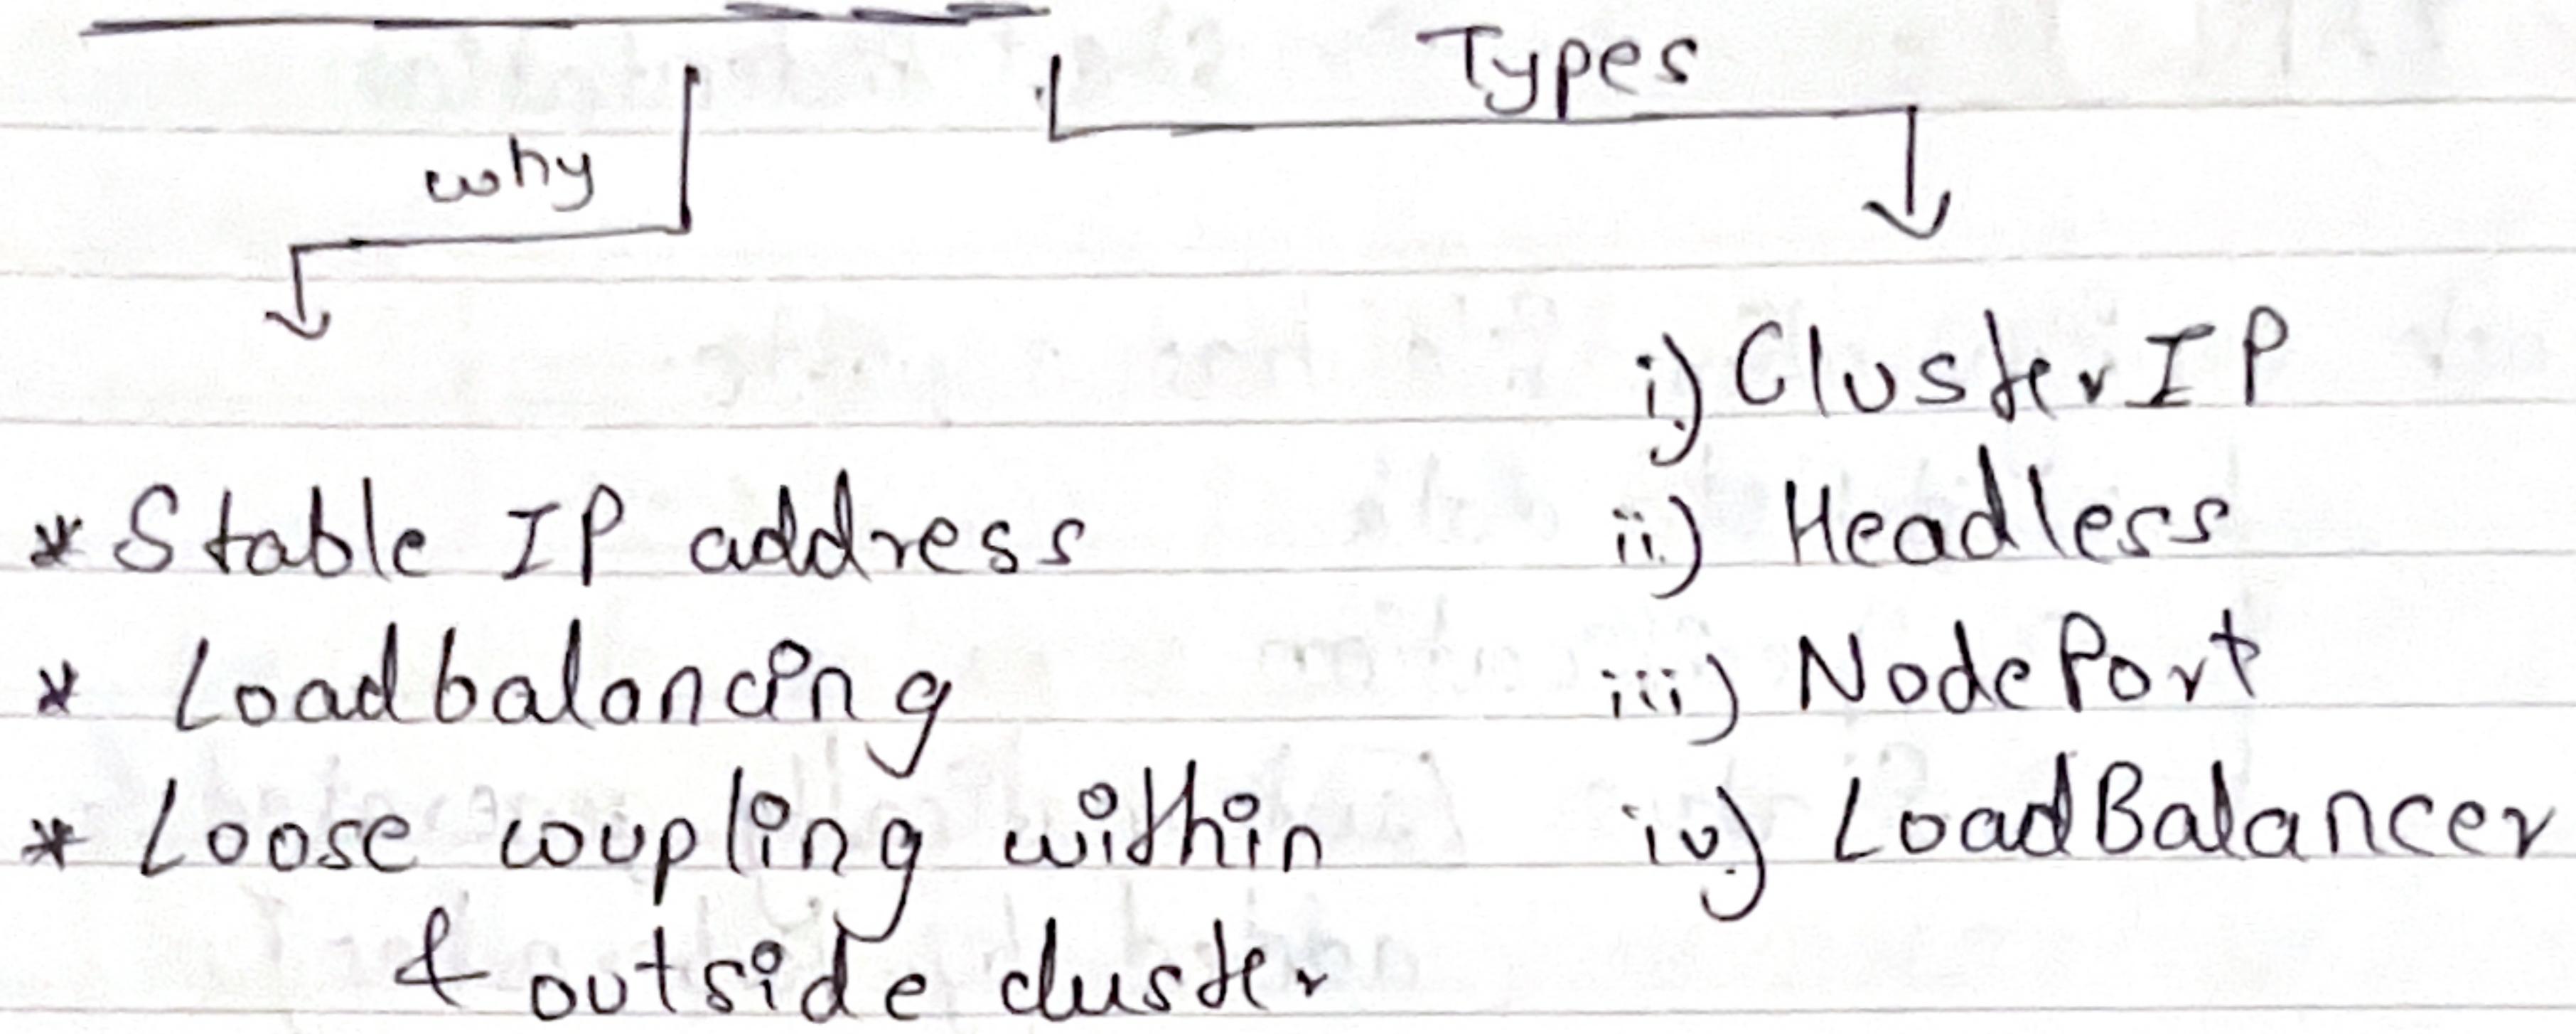
Dictionary

↳ Unordered

List

↳ Ordered

Kubernetes services



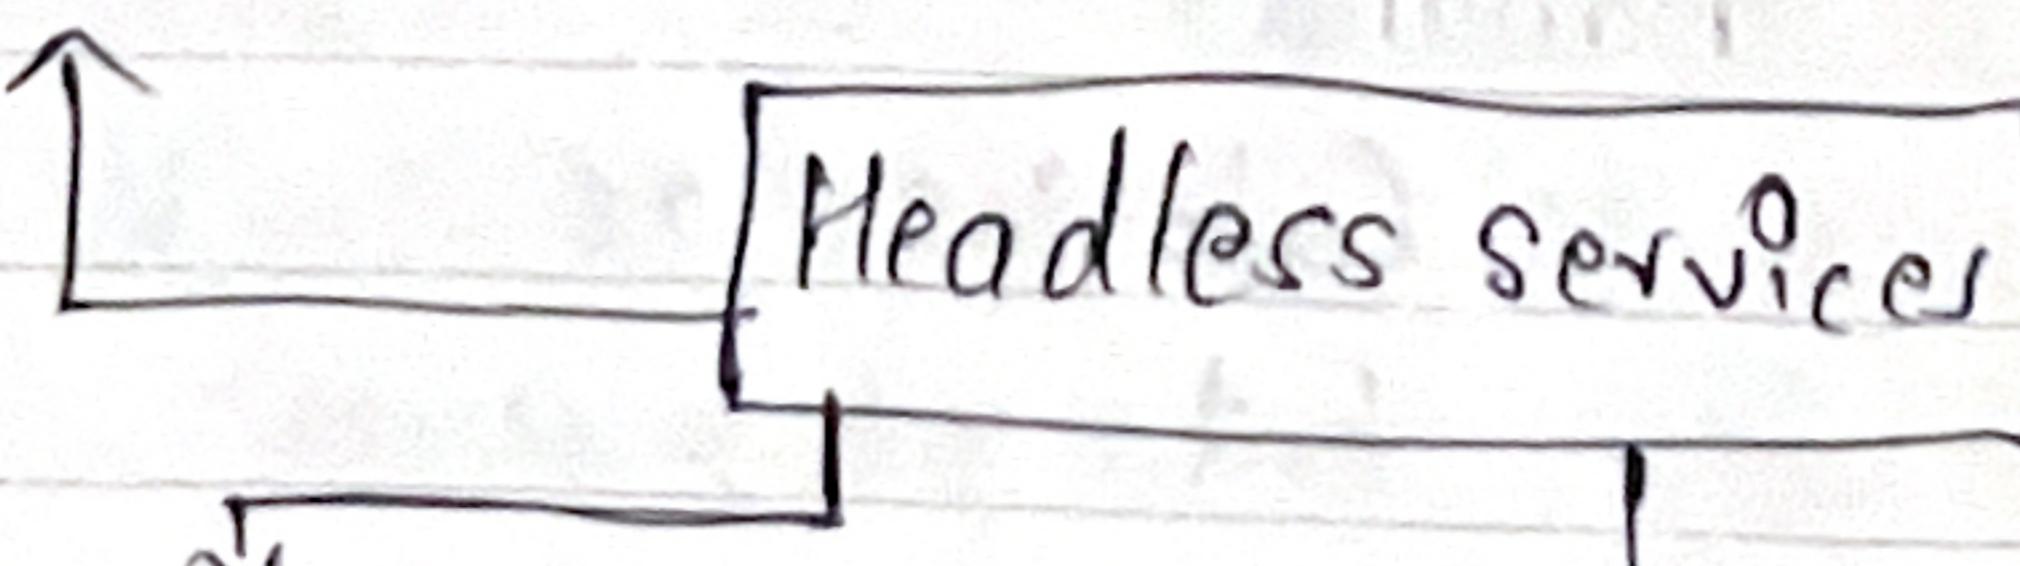
ClusterIP services

Default Type

- * Pods are identified via selectors
- * Service Port is arbitrary
- * Target port must match port the container is listening at

Client wants to communicate with 1 specified pod directly

Pods want to talk directly with specified pod

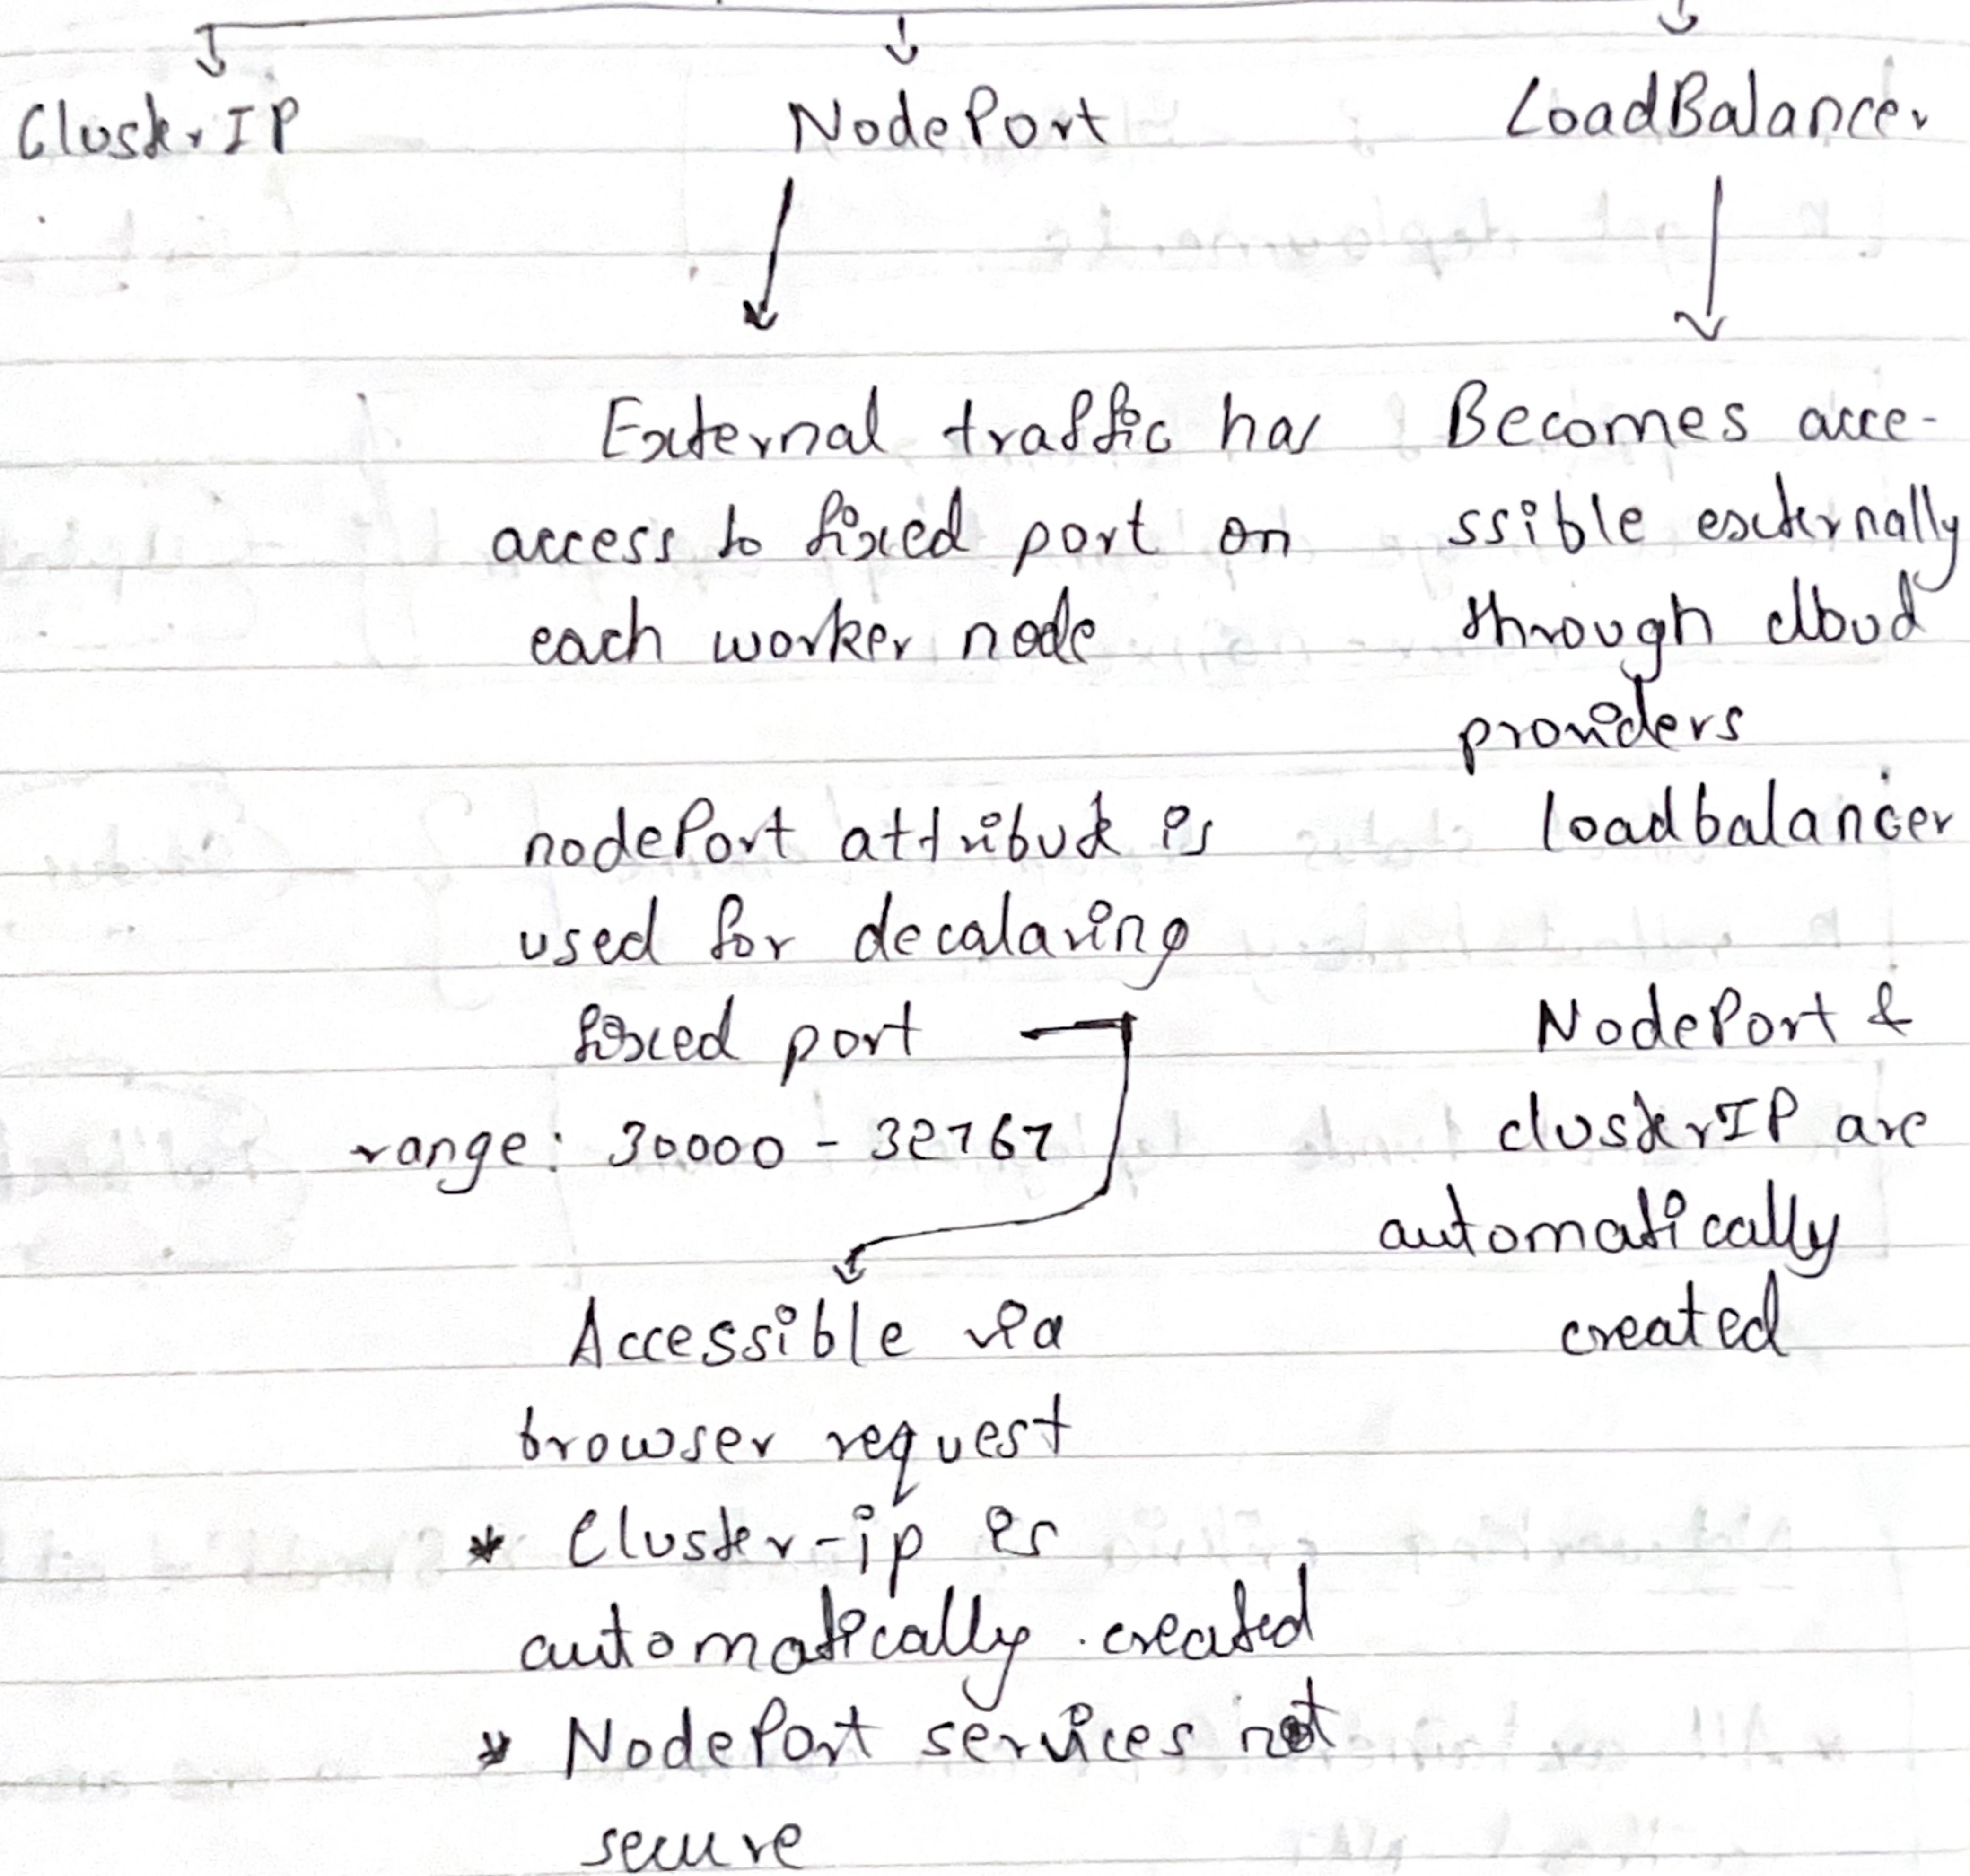


Use case:

Stateful applications like databases

→ We can set Cluster IP to none to use headless service

3 Service type attributes

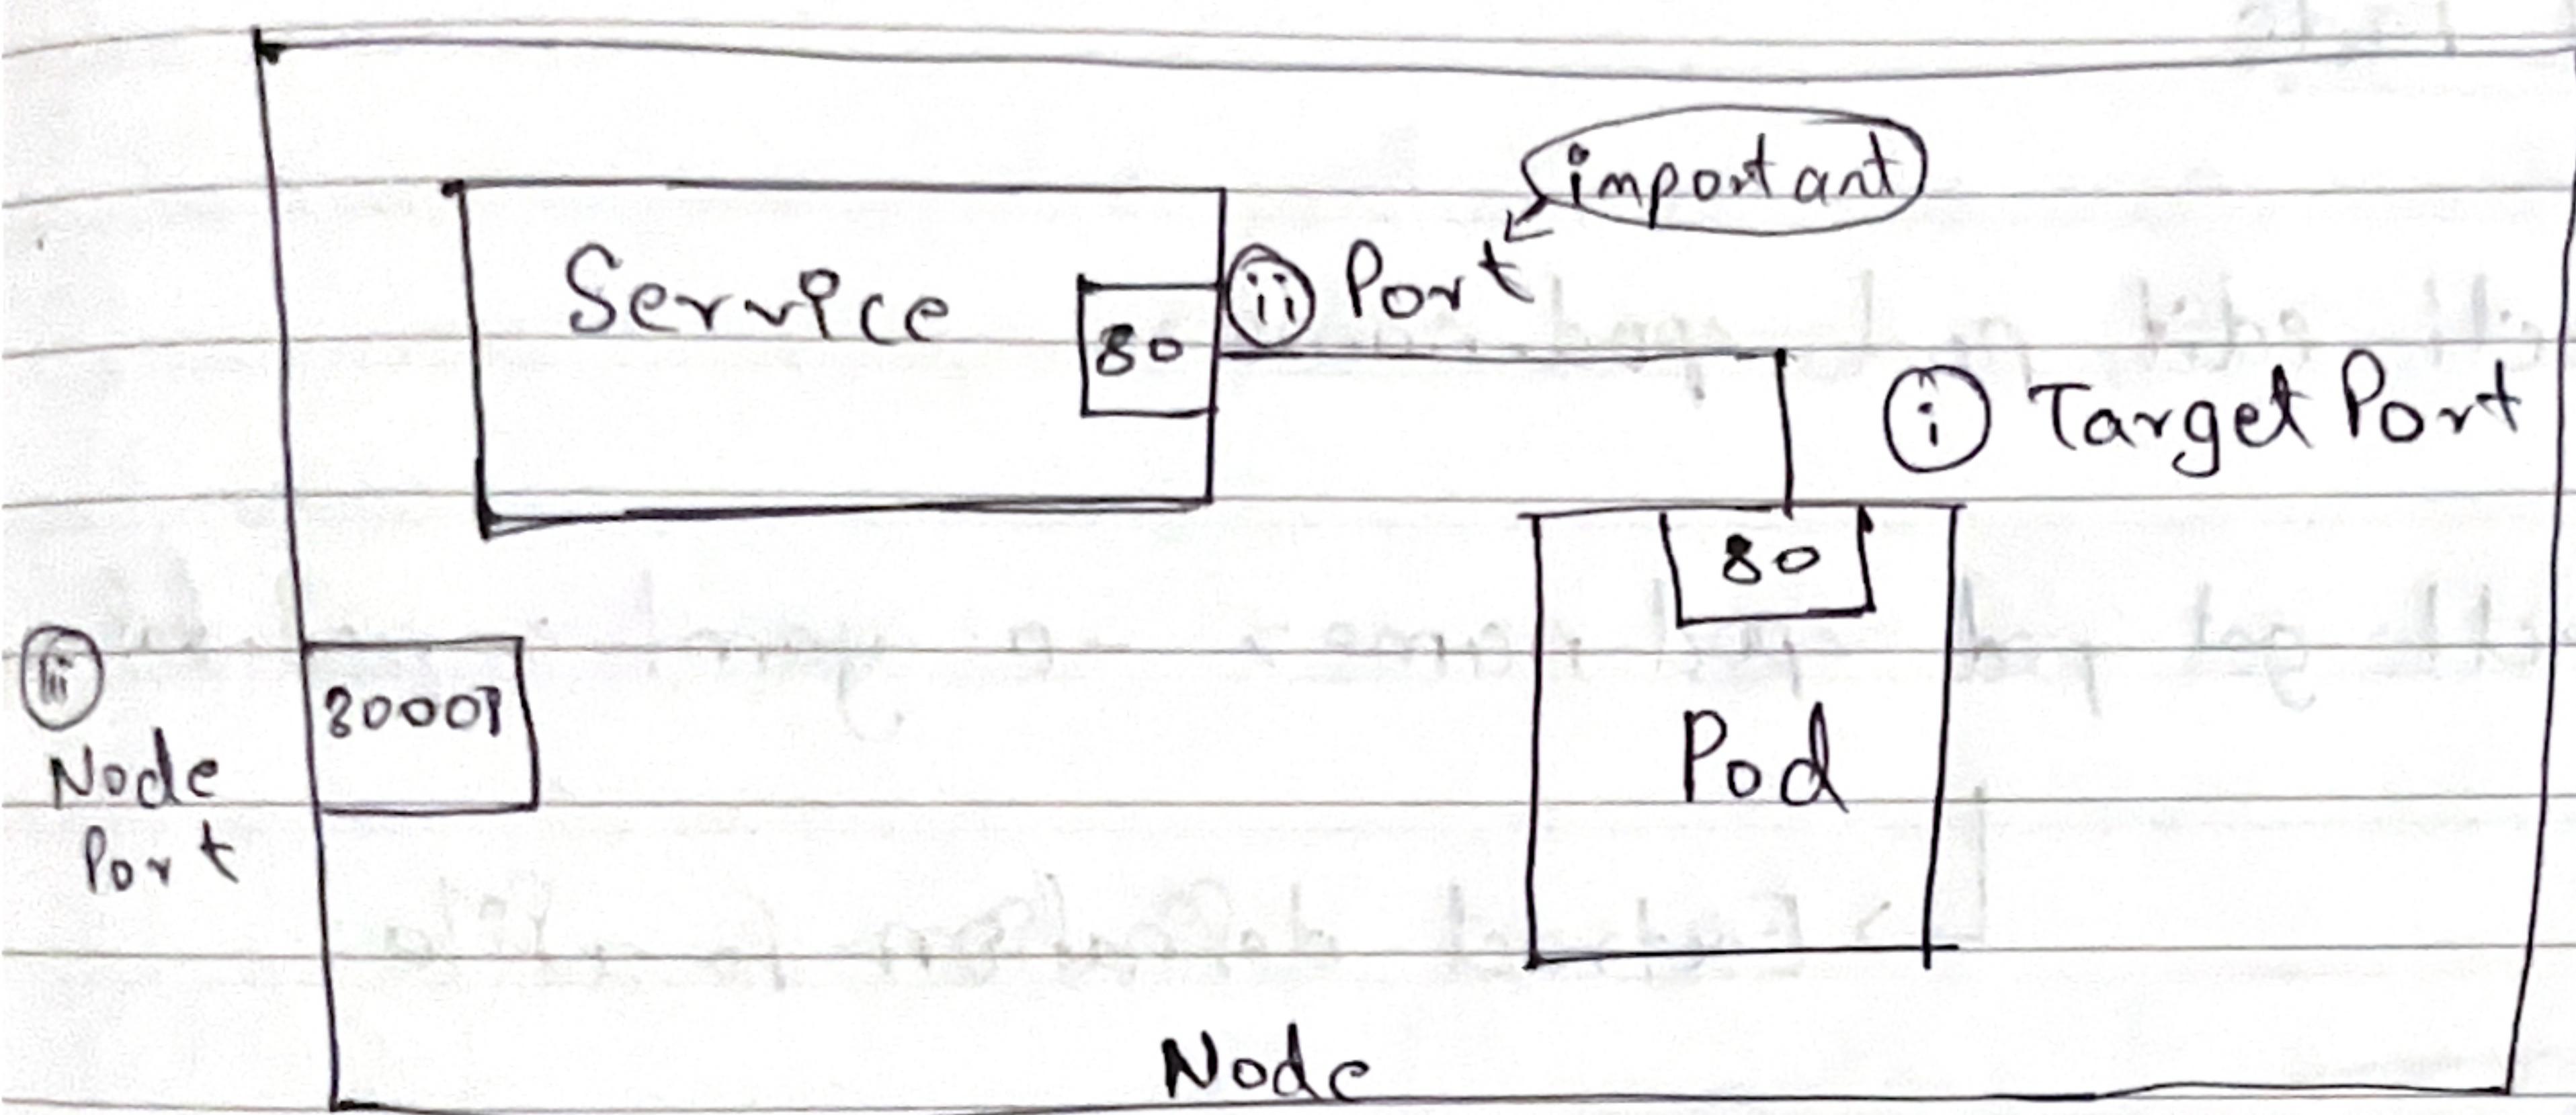


* LoadBalancer service is extension of NodePort

* NodePort is extension of ClusterIP

* Configure Ingress or LoadBalancer for production environments

Service - NodePort



{
 `svc-def.yaml`}

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-svc
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80*
      nodePort: 30008
  selector:
    app: myapp
    type: Frontend
```

{
 `pod-def.yaml`}

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: Frontend
spec:
  containers:
    - name:
      image:
```

* Not Required

Update & Rollback deployment commands

```
kubectl create -f <fileName>  
k get deployments
```

- Create
- Get

```
kubectl apply -f <fileName>  
kubectl set image deployment/app-deployment  
nginx=nginx:1.9.1
```

- Update

```
kubectl rollout status deployment <name>  
kubectl rollout history <name>
```

- Status

```
kubectl rollout undo deployment <name>
```

- Rollback

Networking criteria in Cluster → Should be set by us

- * All containers / PODs can communicate to one another without NAT
- * All nodes can communicate with all containers & vice-versa without NAT
- We can check status of each revision individually by using `--revision=<>/flag`
- * We can use `--record` flag to save command against the revision number
- * To rollback to specific version we will use `--to-revision = <revision number>`